

The LSU logo is displayed in a bold, purple, sans-serif font. It is positioned in the upper left corner of the slide, overlaid on a semi-transparent image of a tree branch and the Audubon Hall building.

**LSU**

School of

Electrical Engineering and Computer Science



06.10.16

# Garbage Collection for General Graphs

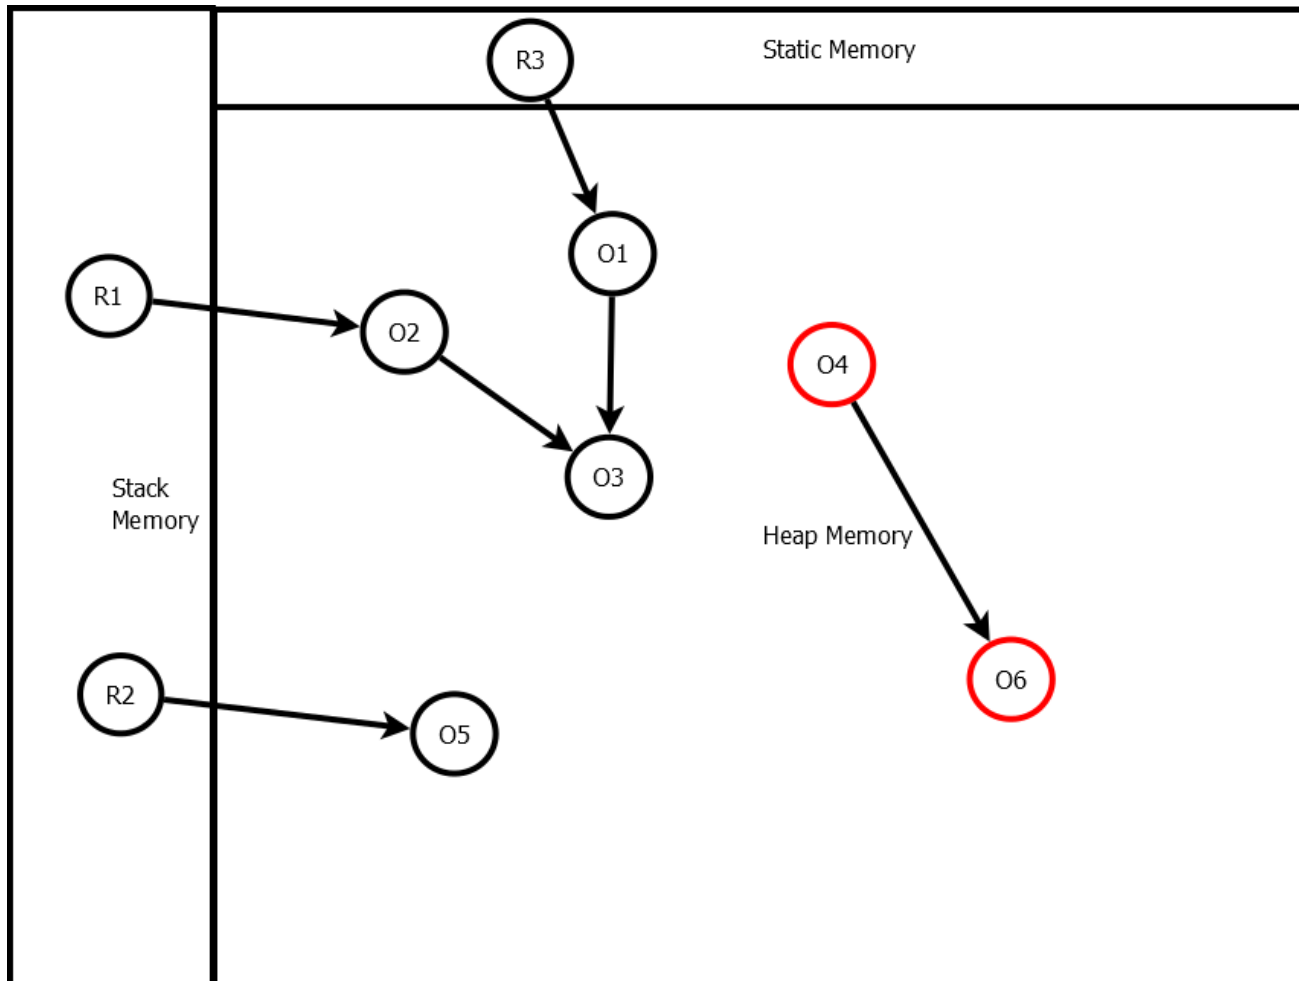
Hari Krishnan

LOVE PURPLE  
LIVE GOLD

# Garbage Collection (GC)

- In heap memory, objects hold pointers/references/links to other objects in the heap.
- Stack variables and static variables hold references to objects in the heap. They are called roots ( R ).
- An object is said to be reachable / live if there is any path from a root to the object.
- Unreachable objects are called garbage and memory allocated for those objects can be reclaimed for future use.
- Garbage collection is the process of collecting unreachable objects in the heap.

# Garbage Collection (GC)







# How to collect garbage?


- Manual Memory Management ( $M^3$ ) and Automatic Memory Management (GC).
- $M^3$  is used by programmers who use languages with no managed run-time systems such as C/C++.
- GC is the thread that runs in the runtime systems(managed runtime systems) to automatically detect the garbage and reclaim them.
- GCs are widely available in various runtime systems including Java Virtual Machine, LISP, and Scheme systems.

# Dangling Pointers


- Dangling pointer is a pointer that points to a deleted object.
- Some other object is allocated in the same address in the interim time.
- In  $M^3$ , simple reference counting based smart pointers can sometimes solve this problem in a concurrent programming environment.
- We focus on concurrent programs and their environments in this presentation.




# Dangling Pointers



Thread 1



Thread 2



# Double-free Bugs (DFB)

- DFB is pointer that points to a deleted object calls delete again.
- Some other object is allocated in the same address in the interim time and creates dangling pointers, or crashes.
- In  $M^3$ , simple reference counting based smart pointers can solve this problem in concurrent programming environment in an acyclic graph.
- Otherwise, timestamps are used in lock-free algorithms to avoid this problem (similar to ABA ).




# Memory Leak


- Objects in heap memory are not referenced by any roots but not deleted.
- Low memory utilization.
- Cycles in the heap cannot be reclaimed by smart pointers.





# Memory Leak



Thread 1




Thread 2




R2 to O1 edge deletion will delete the cycle due to weak reference from O3 to O1.



# Memory Leak



Thread 1




Thread 2



R1 to O2 must be deleted before R2 to O1 to avoid dangling pointers



# Memory Leak



Thread 1




Thread 2




R2 to O1 must be deleted before R1 to O2 to avoid dangling pointers





# Dangling Pointer



Thread 1



Thread 2



Uncertainty in order of operations make it difficult to use the smart pointers efficiently.

# Advantages of GC

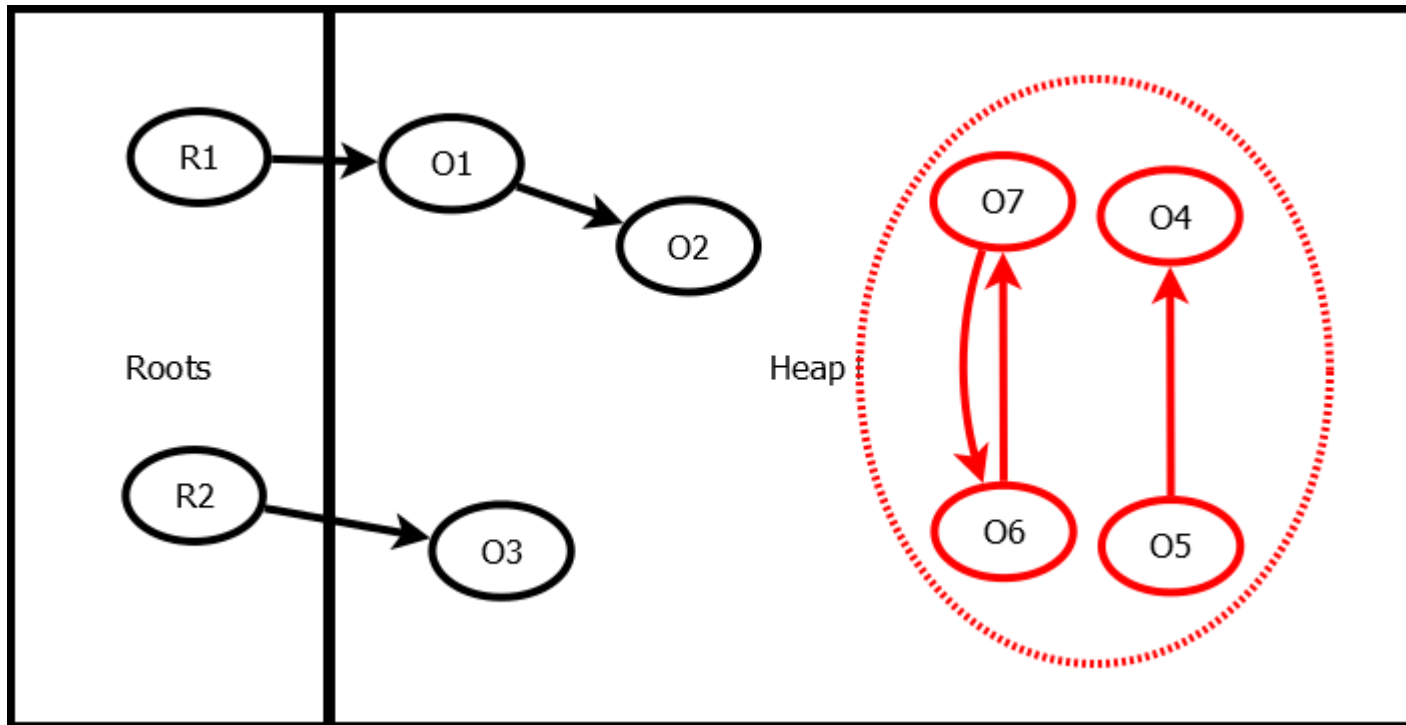
- GC has the global knowledge solves the problem of dangling pointers, double-free bugs and memory leaks.
- Reduces number of lines of code to write.
- AHA! Boehm and Spertus announced that in the next C++ standard, GC can be expected!
- Commercial Software Development research claims GC reduces development cost. [Butters, 2007]
- Useful for distributed object stores, parallel and distributed programming languages, distributed databases, and WWW.

# Mark-Sweep

- When the amount of memory consumed reached its threshold, the collector starts marking the nodes reachable from the roots.
- ( + ) Mark and Sweep only needs one bit and can be easily made concurrent.
- ( - ) Garbage cannot be collected promptly.
- ( - ) The whole allocated heap is traversed for each collection as they first traverse all the live nodes and then they traverse all the dead nodes to delete.




# Mark-Sweep



# Reference Counting (RC)

- Each object has a RC that denotes how many incoming references it has.
- When an object has zero RC, it is garbage.
- ( - )The method cannot delete cyclic garbage as all cyclic garbage nodes have positive RC.
- ( + )Decision are made locally.
- ( + ) The entire heap is not traversed to delete garbage nodes.

# Reference Counting





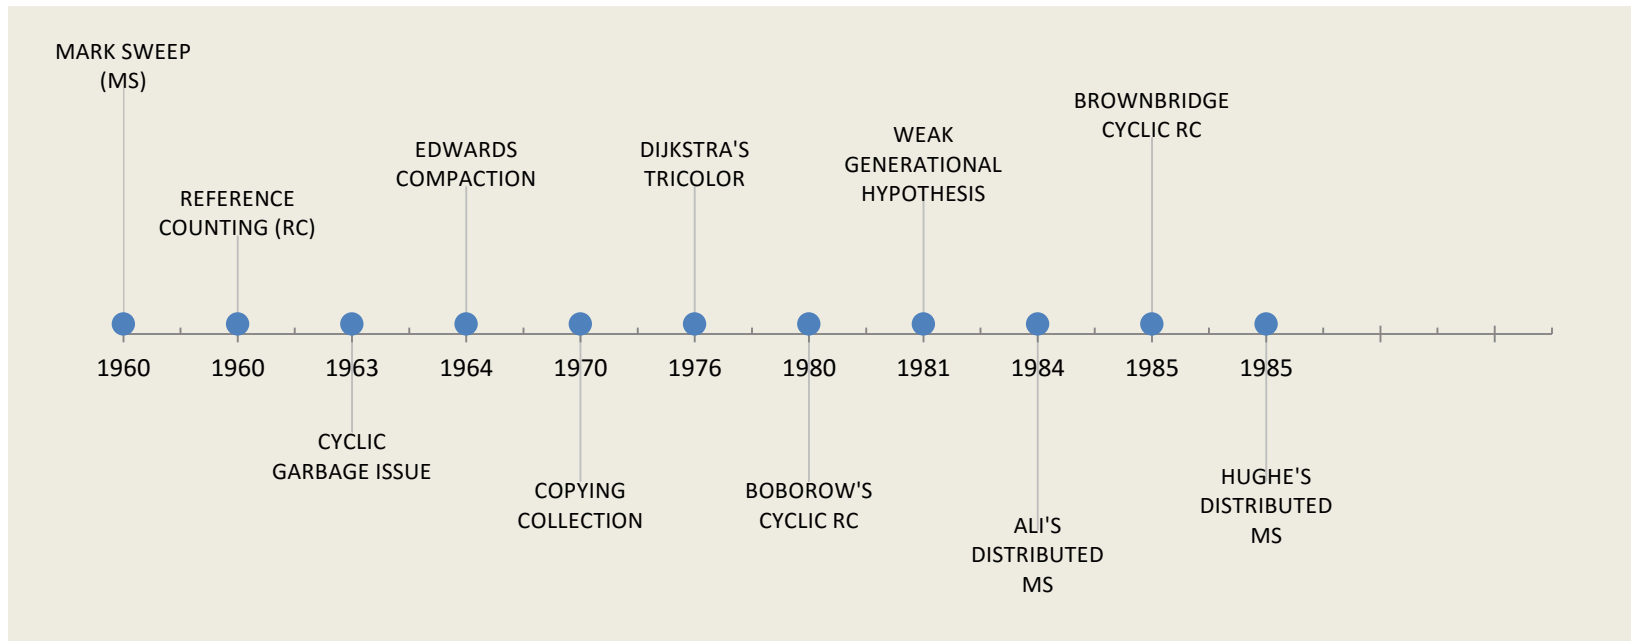
# GC Desired attributes

- Concurrent GC (Low Pause Time)
- Multi-collector GC (Parallel and High Throughput)
- No global Synchronization (High Throughput)
- Locality-based GC (High Throughput)
- Scalable
- Prompt
- Safe
- Complete

# Literature Review

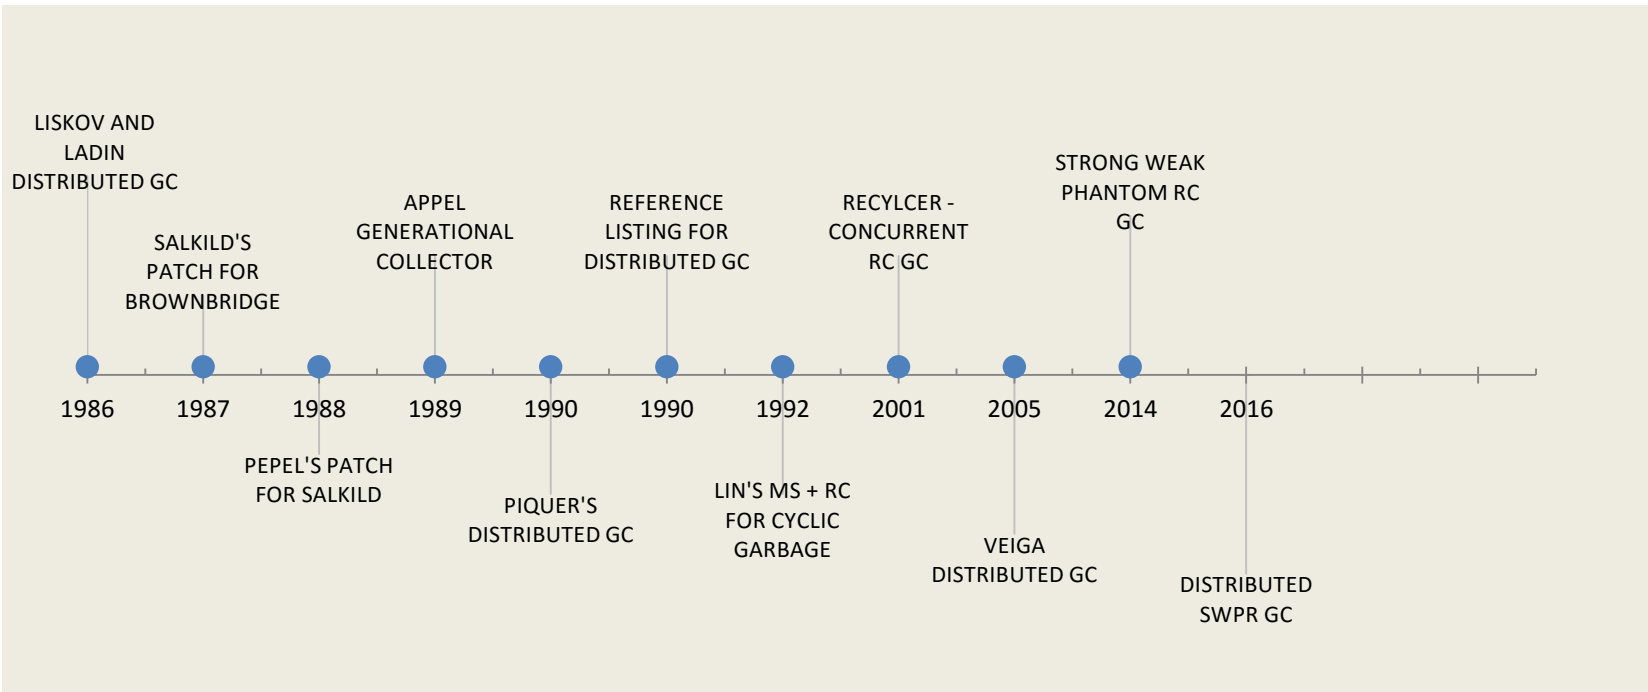
Name	C	MCA	NGS	Locality	Scalable	Safe	Complete	Prompt	S/D
Mark Sweep	Yes	Yes	No	No	No	Yes	Yes	No	S & D
Reference Counting	Yes	No	Yes	Yes	No	Yes	No	Yes	S
Cyclic Reference Counting	Yes	No	Yes	Yes	No	Yes	Yes	Yes	S
Generational	Yes	No	No	No	No	Yes	Yes	No	S
Liskov	No	No	No	No	No	Yes	Yes	No	D
SWP (Thesis Algo – 1)	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	S
SWPR (Thesis Algo – 2)	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	D

# Literature Review





# Literature Review



# Slides Organization


- Preliminaries – Brownbridge Method
- SWP – (Shared Memory GC)
- SWPR (Distributed Memory GC)

# Brownbridge Method


- RC and tracing technique to collect cyclic garbage.
- It uses two reference counts: Strong Reference Count (SRC) and Weak Reference Count (WRC).
- All live nodes have positive SRC.
- No strong cycles are allowed.
- When a reference is created to a node, if the target node already has outgoing references, it is considered a weak reference.




# Brownbridge Method



# Recover Case



# Premature Delete Case





# Salkild's and Pepel's Work

- Salkild eliminated the premature deletion but introduced non-termination in certain cases.
- Pepel improved Salkild's work with the trade-off of exponential cleanup cost.
- Practically all of the attempts to correct the Brownbridge failed.

# Shared Memory GC


# SWP

- Strong-Weak-Phantom(SWP) Reference Count Shared Memory GC.
- The idea is to create a path from a root to each live node through strong references and avoid creating cycles of strong reference by using weak references.
- **Strong Cycle Invariant** : No strong cycles will exist in the reference graph at any point in time.
- **Weak Heuristic** : All weak edges are not cycle closing edges.
- **Edge Label Heuristic**: An edge will be weak when the target node has outgoing edges at the time of creation.
- Phantom is a transient state that is used only in the cycle detection algorithm.




# SWP

Counters	State
$SRC > 0$	Live
$SRC = 0 \ \& \ WRC > 0$	Potential Cyclic Garbage
$PRC > 0$	GC processing node
$SRC = 0 \ \& \ WRC = 0 \ \& \ PRC = 0$	Garbage



# SWP & Header

- Apart from reference counts, each node also has a which bit, phantomized flag, and an array of which bits for outgoing references.
- Process lists and request queues are used.



# Delete Edge

When a node loses its last strong reference:

```
if(WRC==0 && PRC==0)
```

```
    start deleting the node and delete all outgoing edges
```

```
else if(WRC>0 )
```

```
    convert all the incoming weak references to strong and phantomize  
    outgoing edges
```

# Three phases

- Phantomize
- Recovery
- Cleanup

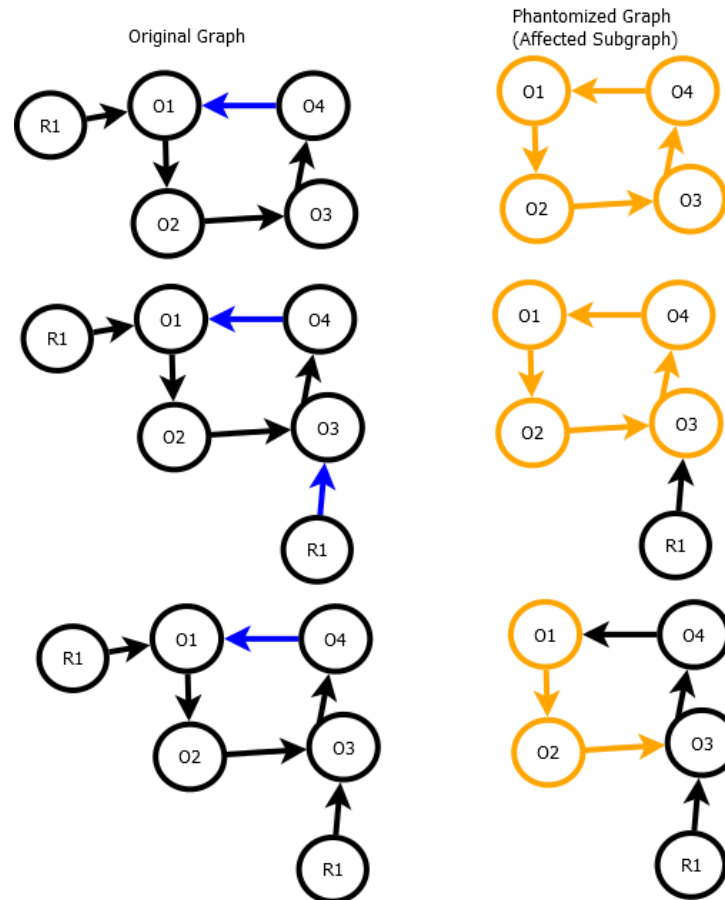
These three phases happen for each traversal initiated. Each collector thread maintains list of nodes it processed in each phase.



# Phantomization

- Decrement all internal references.
- Phantomization stops when it reaches a phantomized node, or a live node.
- Nodes toggle if any weak reference remain after phantomization.
- Phantomization propagates when any of the stop criteria is not met.

# Phantomization




# Recovery


- If any external reference found, correct the graph to strong / weak graph
- Delete nodes from process list as they are recovered.
- Process starts by scanning all the process list contents and verifying if any of them has any strong references after phantomization.

# Recovery


Original Graph



Phantomized Graph



Recovered Graph





# Concurrent and parallel Collector

- Atomic operations can be used to access RCs.
- SWP collection does not stop the application.
- Requests are queued.
- Collector thread processes the queue and starts the requested tasks.
- Collectors write their own id in collector id in parallel collector mode.
- When the collector visits a node with a different id, then they synchronize.

# Concurrent and parallel Collector

- **No Partition Principle** : When two or more collectors synchronize, they merges their lists of processed nodes and one of the collectors takes over the processing of those nodes.
- Parallel collector adds one more attribute to the object header.

# How SWP can be used?


- The SWP algorithm is directly applicable to shared memory systems.
- The SWP algorithm is also applicable to distributed system with centralized queues.
- Unlike other distributed GC systems which require the application to be halted, this one does not require it.
- Mobile actor based collectors do not require centralized queues and are directly applicable to the distributed systems with message size proportional to graph size.

# Distributed Memory GC



# SWPR GC

- SWPR - > (Strong, Weak, Phantom, Recovery).
- Distributed Algorithm.
- Distributed Termination Detection is used. (Parent attribute and Wait Count).




# Distributed Model

- Congest:  $O(\log n)$  message size.
- No reference listings allowed.
- Nodes can send messages only to neighbors.
- Asynchronous network model.

# Node Classification


- A -> Initiator
- B, B' -> Purely Dependent Set
- C, C' -> Partially Dependent Set
- D, D', E, E', F -> Independent Set
- C', D', E', F -> Supporting Set
- F, D', E' -> Build Set
- C' -> Recovery Set





# Node Classification


- A -> Initiator
- B, B' -> Purely Dependent Set
- C, C' -> Partially Dependent Set
- D, D', E, E', F -> Independent Set
- C', D', E', F -> Supporting Set
- F, D', E' -> Build Set
- C' -> Recovery Set






# Node Classification

- A -> Initiator
- B, B' -> Purely Dependent Set
- C, C' -> Partially Dependent Set
- D, D', E, E', F -> Independent Set
- C', D', E', F -> Supporting Set
- F, D', E' -> Build Set
- C' -> Recovery Set




# Node Classification

- A -> Initiator
- B, B' -> Purely Dependent Set
- C, C' -> Partially Dependent Set
- D, D', E, E', F -> Independent Set
- C', D', E', F -> Supporting Set
- F, D', E' -> Build Set
- C' -> Recovery Set



# Node Classification


- A -> Initiator
- B, B' -> Purely Dependent Set
- C, C' -> Partially Dependent Set
- D, D', E, E', F -> Independent Set
- C', D', E', F -> Supporting Set
- F, D', E' -> Build Set
- C' -> Recovery Set





# Node Classification


- A -> Initiator
- B, B' -> Purely Dependent Set
- C, C' -> Partially Dependent Set
- D, D', E, E', F -> Independent Set
- C', D', E', F -> Supporting Set
- F, D', E' -> Build Set
- C' -> Recovery Set






# Node Classification

- A -> Initiator
- B, B' -> Purely Dependent Set
- C, C' -> Partially Dependent Set
- D, D', E, E', F -> Independent Set
- C', D', E', F -> Supporting Set
- F, D', E' -> Build Set
- C' -> Recovery Set



# State Diagram



# Phantomization


- Converts all the internal edges in the non-independent nodes of the subgraph to phantom.
- The process makes sure the internal edges are not counted for the decision process.
- External weak edges are converted into strong during this process.
- The process uses a forward phase and a backward phase to finish operations.



# Phantomization


Original Graph

Initiator




Forward Phase Phantomized Graph

Initiator



Reverse Phase Phantomized Graph

Initiator

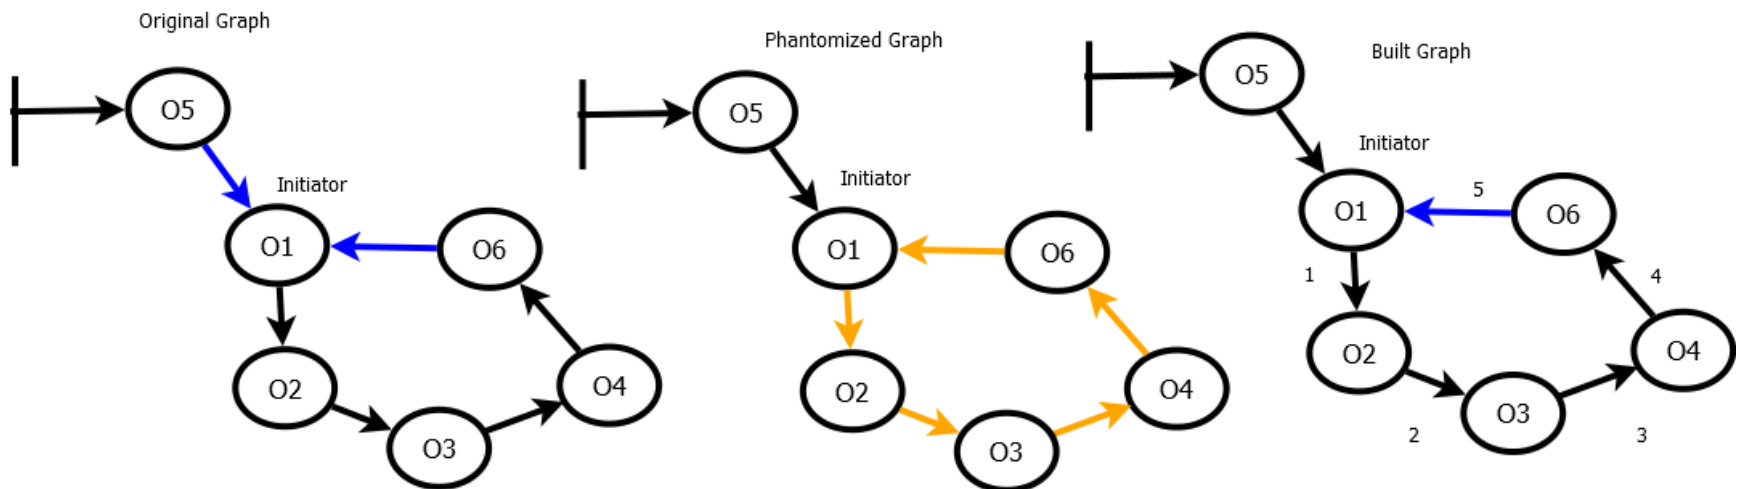




# Correction

- Recovery, Build, and Plague Delete are correction phases.
- After phantomization, if the initiator has nodes in the build set, it converts all the phantom edges in the subgraph into strong / weak based on weak heuristic.
- Test to find build set: true If the initiator has any strong edges after phatonmization.
- Build by initiator transforms the graph back to regularity.

# Build Phase




# Recovery


- Recovery is a complicated phase.
- If initiator does not have any build set, recovery messages are sent.
- A recovery message will only affect the subgraph that is not yet processed.
- On the reverse phase, a recovery node can start building too.

# Recovery


Original Graph




Phantomized Graph



Forwarded Phase Correction (Recovery)



Reverse Phase Correction (Building)






# Delete


- A node will be deleted only if all the counts are zero.
- A node on receiving delete will send delete only if there is no strong incoming edge.
- Delete is invoked by the initiator if there is no build set and no recovery set.

# Delete


Original Graph




Phantomized Graph




Recovery failed Graph




Delete Initiated



O2 is deleted



O3 is deleted



# Isolation property

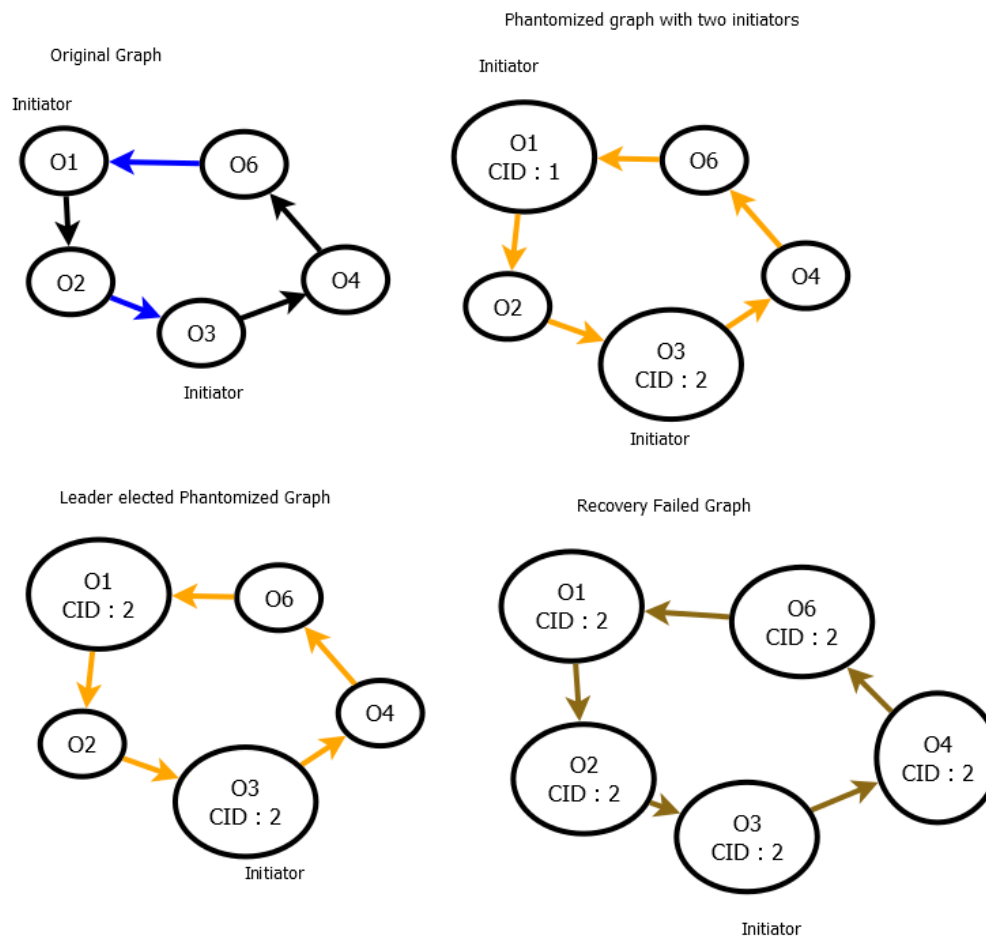
- The affected subgraph is not mutated.
- Mutation to the affected subgraph occurs, but do not affect the decisions of the initiator to delete or build.
- Do not affect the decision of recovery set to build by the correction phases.

# Symmetry Breaking

- When multiple collector process the same subgraph, there is possibility of cycle dependency.
- A leader needs to be elected among conflicting collector operations.
- Higher collection id is chosen as leader when conflict happens.
- Each node contains a collector id in the correction phase.
- All nodes prefer to be in a collection with higher id.
- All collection ids are unique and so there is a total ordering among collections.



# Acyclic Principle (Isolation)

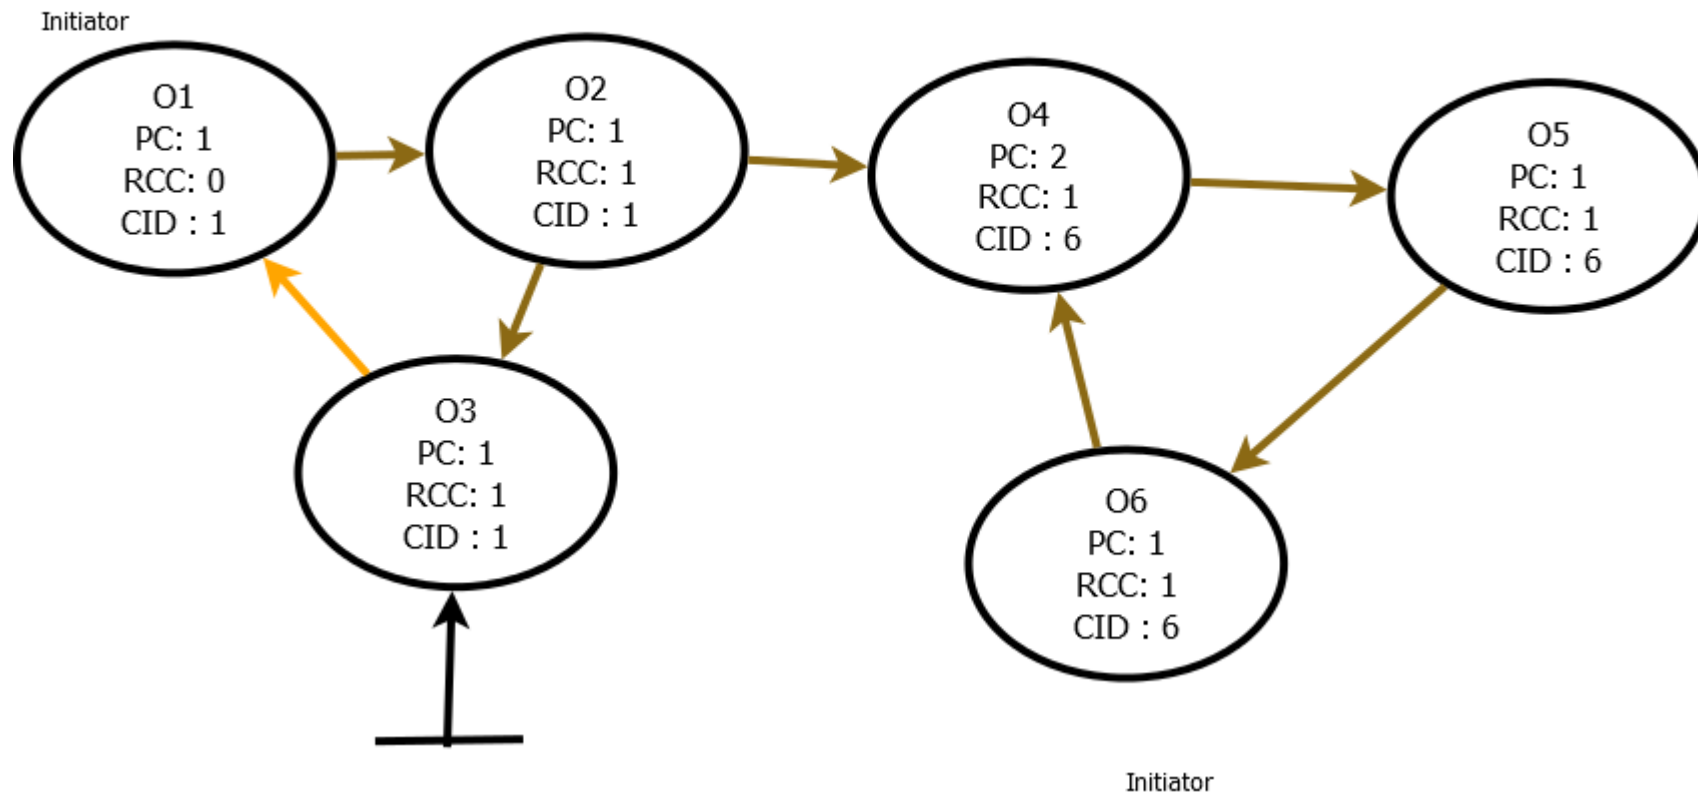


# Recovery Count

- For the recovery phase to start reverse phase, it must have recovery count equal to phantom count.
- Every recovery message received increments recovery count.
- This creates ordering among subgraphs based on topology, regardless of uncertainty in the collection ids.
- Low collector id cannot increment higher collection recovery count.

# Topology Ordering (Isolation)

CID 6 is in forward recovery phase and CID 1 is also in forward recovery phase and topology dictates ordering of collection using Recovery Count (RCC).






# Transactional Approach

- When two collectors' operations meet at a point, the lower collection id will not proceed because of symmetry breaking rules.
- To complete the computation of the lower collection process during recovery, recovery phases alone are restarted. Other phases work seamlessly with multiple collectors colliding.
- When multiple collectors of different phases meet, we wait until the reverse phase finishes and then redo if they need to upgrade.




# Redo transaction (Isolation)


Original Graph




Phantomized Graph due to O3



O3 is building but O7 creates Phantomization




O3 joins O7 and redo phantomization of O3 phantomized graph



# Recovery Start Over

Small Collector in forward recovery phase conflicting with big recovery collection must restart recovery phase again.  
CID 1 will restart recovery due to collision with CID 6 recovery.

Initiator



Initiator

# Complexities

- Finishes in  $O(E)$  time without conflict in asynchronous systems.
- All messages contain only id of the node ( $O(\log n)$  size), collector id, constants such as type of message, and some flags. So  $O(\log(n))$  message size.
- All nodes requires constant space to save the collector id, parent, and flags.

# Conclusion & Future Work

- Both of our algorithms satisfy all the desired properties of the GC algorithms.
- Most run-time systems cannot guarantee the order of destructors being called. Our algorithm is not designed to solve the order of destructor problem.
- In distributed systems, space, network and other failures are unavoidable. Future directions for this thesis would be designing an algorithm to incorporate fault-tolerance.